

Preview

- ❑ System Calls vs. Library Functions
- ❑ File Descriptors for a process
- ❑ System Call for Managing Files
 - write()
 - read()
 - open()
 - close()
 - lseek()
 - pwrite(), pread();

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

1


System Calls vs. Library Functions

- ❑ A **system call** is a request for the operating system to do something on behalf of the user's program.
- ❑ The system calls are functions used in the kernel itself.
- ❑ To the programmer, the system call appears as a normal C function call.
- ❑ When a system call, control change from user's mode to kernel's mode.

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

2

System Calls vs. Library Functions



COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

3

System Calls vs. Library Functions

- ❑ System calls is the layer of software interface to the kernel.
- ❑ It is controlled by kernel (unbuffered I/O)
- ❑ Library functions are built on top of system call (Buffered I/O). It is called and controlled by a process.
- ❑ Linux system calls are used to process management, file system management, and inter-process communication.
- ❑ The Linux system interface consist of about 80 system calls.
- ❑ To access and control file, need system call to open, read, write, close file.

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

4

System Calls vs. Library Functions

- ❑ A system call is very expensive routine
 - Change a mode from user's mode to kernel's mode.
 - Save all parameters for user's program for later execution (save snapshot of CPU)
 - Load all necessary parameters for system call routine to CPU .
 - Execute the system call routine.
 - After system call routine, load all saved parameters for user's program to CPU.
 - Change mode from kernel's mode to user's mode.
 - Continue execution of user's program.

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

5

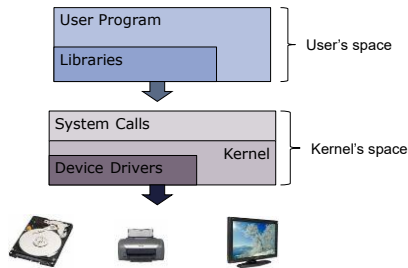
System Calls vs. Library Functions

- ❑ To reduce the overhead of system calls, system such as Unix or Linux provide library functions use buffers.
- ❑ For example, I/O function library that provide buffered output.
- ❑ With these library functions, system call routine need execute when size of buffer becomes full.
- ❑ This dramatically reduce the system call overhead.

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

6

System Calls vs. Library Functions



COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

7

File Descriptors for a Process

- ❑ To the kernel, all open files are referred to by file descriptors.
- ❑ A file descriptor is a non-negative integer.
- ❑ When we open an existing file or create a new file, the kernel returns a file descriptor to the process.
- ❑ We can read or write a file with the file descriptor which was return by system calls `open()` or `create()`.

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

8

File Descriptors for a Process

- ❑ A process has a number of the descriptors associated it.
- ❑ When a process is created, it has three descriptors which can be used for input, output and error.
 - 0: Standard input (from keyboard by default)
 - 1: Standard output (to screen by default)
 - 2: Standard error (to screen by default)

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

9

System Calls for Managing Files (`write()`)

- ❑ The `write()` system call attempts to write *nbyte* bytes from the buffer pointed to by *buf* to the file associated with the open file descriptor, *filides*.
- ❑ It returns the number of bytes actually write.

```
#include <unistd.h>
ssize_t write(int filides, const void *buf, size_t nbyte);
```

Returns: number of bytes written, or -1 on error

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

10

System Calls for Managing Files (`write()`)

```
//write.c
#include <unistd.h>
#include <stdlib.h>

int main()
{
    if ((write(1, "Hear is some data\n", 18)) != 18)
        write(2, "error on file descriptor 1\n", 46);
    exit(0);
}
```

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

11

System Calls for Managing Files (`write()`)

```
#include <unistd.h>
#include <stdlib.h>

int main()
{
    char Buffer[]="Hear is some data\n";
    if ((write(1, Buffer, 18)) != 18)
        write(2, "error on file descriptor 1\n", 46);
    exit(0);
}
```

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

12

System Calls for Managing Files (read())

- The `read()` system call function attempts to read *nbyte* bytes from the buffer pointed to by *buf* to the file associated with the open file descriptor, *fildev*.
- It returns the number of bytes actually read.

```
#include <unistd.h>
ssize_t read(int fildev, const void *buf, size_t nbyte);
```

Returns: size of byte read, 0 if end of file, -1 on error

System Calls for Managing Files (read())

```
/* read_buffer.c
 * prepare 128 bytes buffer and read a string upto 128 bytes*/
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>

int main()
{
    char buffer[128];
    int nread;
    /* can read up to 128 character from stdin */
    nread = read(0, buffer, 128);

    if (nread == -1)
        write(2, "A read error\n", 12);
    /* write content of buffer to stdout */
    if ((write(1, buffer, nread)) != nread)
        write(2, "A write Error\n", 14);
    exit(0);
}
```

System Calls for Managing Files (read())

```
/* read_buffer1.c */
/* read and write byte by byte until empty (Ctrl-D)*/
#include <stdio.h>
#include <unistd.h>

int main()
{
    char b[1];
    int nread;
    /* read character by character from stdin */
    while ((nread = read(0, b, 1) > 0))
        write(1, b, nread); /* write char by char to stdout */

    return 0;
}
```

System Calls for Managing Files

- Only using `read` and `write` system call, we can copy standard input to standard output.
- This assume that these have been set up by the shell before this program is executed.

System Calls for Managing Files

```
/* copystdio.c copy standard input to standard output */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h> /*STDIN_FILENO, STDOUT_FILENO*/

/* buffer size effect the efficiency of the program */
#define BUFFER_SIZE 2
void err_sys(char *str)
{
    printf("%s", str);
    exit(1);
}

int main()
{
    int nbyte;
    char buffer[BUFFER_SIZE];

    while ((nbyte = read(STDIN_FILENO, buffer, BUFFER_SIZE)) > 0)
        if (write(STDOUT_FILENO, buffer, nbyte) != nbyte)
            err_sys("Write Error");
    if (nbyte < 0)
        err_sys("read Error");
    exit(0);
}
```

System Calls for Managing Files (open())

- `open()` lets you open a file for reading, writing, or reading and writing.
- It returns a file descriptor.
- The prototype for the `open()` system call is:

```
#include <fcntl.h>
int open(const char *fname, int flags)
int open(const char *fname, int flags, mode_t mode);
```

Returns: file descriptor or -1 on error

- The third argument mode is used only when a new file is being created.

System Calls for Managing Files (open())

- The allowable flags as defined in "/usr/include/fcntl.h" are:

```
#define O_RDONLY 0    /* Open the file for reading only */
#define O_WRONLY 1    /* Open the file for writing only */
#define O_RDWR 2     /* Open the file for both reading and writing */
#define O_NDELAY 04   /* Non-blocking I/O */
#define O_APPEND 010 /* append (writes guaranteed at the end) */
#define O_CREAT 00400 /* open with file create (uses third open arg) */
#define O_TRUNC 01000 /* open with truncation */
#define O_EXCL 02000 /* exclusive open */
```

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

19

System Calls for Managing Files (open())

- The allowable mode_type as defined in a header sys/stat.h" are:

- S_IRUSR Read permission, owner.
- S_IWUSR Write permission, owner.
- S_IXUSR Execute/search permission, owner.
- S_IRGRP Read permission, group.
- S_IWGRP Write permission, group.
- S_IXGRP Execute/search permission, group.
- S_IROTH Read permission, others.
- S_IWOTH Write permission, others.
- S_IXOTH Execute/search permission, others.

Or we can use octal number numerical form

- 0777, 0755, 0555,

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

20

```
/* open.c demonstrate open file */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <fcntl.h> /* defines options flags */
static char message[] = "Hello, world";
#define FILE_MODE (S_IRUSR | S_IWUSR | S_IRGRP | S_IWGRP | S_IROTH | S_IWOTH)

int main()
{
    int fd;
    char buffer[80];
    /* open for rw and create exclusive open, create as rw-rw-rw */
    umask(0); /* clear the mask to be able to create a file with a mode */
    fd = open("datafile.dat", O_RDWR | O_CREAT | O_EXCL, FILE_MODE);
    printf("fd = %d\n", fd);
    if (fd != -1)
    {
        printf("datafile.dat opened for read/write access\n");
        write(fd, message, sizeof(message));
        lseek(fd, 0, SEEK_SET); /* go back to the beginning of the file */
        if (read(fd, buffer, sizeof(message)) == sizeof(message))
            printf("\'%s\' was written to datafile.dat\n", buffer);
        else
            printf("**** error reading datafile.dat ***\n");
        close(fd);
    }
    else
        printf("**** datafile.dat already exists ***\n");
    return 0;
}
```

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

21

```
/* open.c demonstrate open file */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <fcntl.h> /* defines options flags */
static char message[] = "Hello, world";

int main()
{
    int fd;
    char buffer[80];
    umask(0); /* clear the mask to be able to create a file with a mode */
    /* open for rw and create exclusive open, create as rw-rw-rw */
    fd = open("datafile.dat", O_RDWR | O_CREAT | O_EXCL, 0666);
    printf("fd = %d\n", fd);
    if (fd != -1)
    {
        printf("datafile.dat opened for read/write access\n");
        write(fd, message, sizeof(message));
        lseek(fd, 0, SEEK_SET); /* go back to the beginning of the file */
        if (read(fd, buffer, sizeof(message)) == sizeof(message))
            printf("\'%s\' was written to datafile.dat\n", buffer);
        else
            printf("**** error reading datafile.dat ***\n");
        close(fd);
    }
    else
        printf("**** datafile.dat already exists ***\n");
    return 0;
}
```

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

22

System Calls for Managing Files (creat())

- A new file can also be created by calling the create system call.
- One deficiency with creat() is that the file is opened only for writing.

```
#include <fcntl.h>
int creat(const char *fname, mode_t mode);
```

Return: file descriptor, or -1 on error

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

23

System Calls for Managing Files (close())

- Any opened file is closed by a system call close.
- Closing a file releases any record that the process may have on the file

```
#include <fcntl.h>
int close (int filedes);
Returns: 0 if OK, -1 on error
```

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

24

System Calls for Managing Files (lseek())

- Every open file has an associated with **current file offset**.
- It is **non negative integer** that measures the number of bytes from the beginning of the file.
- When a file is opened, offset is set to 0.
- Read/write operation start at current file offset and cause the offset to be incremented by the number of bytes read or write.

System Calls for Managing Files (lseek())

- An open file offset can be explicitly positioned by calling lseek system call.

```
#include <unistd.h>
off_t lseek(int filedes, off_t offset int whence);
```

Returns: new file offset or -1 on error

System Calls for Managing Files (lseek())

- The interpretation of the **offset** depends on the value of the **whence** argument.
 - SEEK_SET**: offset is set to offset bytes from the beginning of the file
 - SEEK_CUR**: offset is set to its current value plus the offset.
 - SEEK_END**: set to the size of the file plus the offset

System Calls for Managing Files (lseek())

```
/* program testlseek.c */
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
#include <sys/stat.h>
#define FILE_MODE (S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH)
/* err_sys display error message and exit */
void err_sys(char *);

int main()
{
    int filedes, offset;

    if ((filedes = open("testlseek.txt", FILE_MODE)) < 0)
        err_sys("Coward File Open Error");

    if ((offset = lseek(filedes, 0, SEEK_END)) == -1)
        err_sys("Coward seek Error");

    printf("offset = %d", offset);
    return 0;
}

void err_sys(char *str)
{
    printf("%s", str);
    exit (1);
}
```

System Calls for Managing Files (lseek())

```
/* lseek.c this program test its standard input to see whether
it is capable of seeking or not */
#include <sys/types.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    int offset;
    if (offset = lseek(STDIN_FILENO, 0, SEEK_END)) == -1)
        printf("cannot seek.\n");
    else
        printf("Seek OK. \n");
    exit (0);
}
```

System Calls for Managing Files (lseek())

```
/* lseek1.c this program display size of input file with input
redirection */
#include <sys/types.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    int offset;
    if (offset = lseek(STDIN_FILENO, 0, SEEK_END)) == -1)
        printf("cannot seek.\n");
    else
        printf("file size is %d bytes.\n", offset);
    exit (0);
}
```

```

/* program createhole.c creat a file with a hole in it */
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
#include <stdlib.h>
#include <sys/stat.h>

char buf1[]="abcdeefghi";
char buf2[]="ABCDEFGHIJ";
void err_sys(char *str)
#define FILE_MODE (S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH)
int main()
{
    int filedes;

    if ((filedes = open("filehole.txt", FILE_MODE)) < 0)
        err_sys("Creat File Open Error");
    if (write (filedes, buf1, 10) != 10)
        err_sys("Creat Write Error"); //now offset = 10*/
    if (lseek (filedes, 40, SEEK_SET) == -1)
        err_sys("Creat seekl Error");//now offset =40 */
    if (write(filedes, buf2, 10) != 10)
        err_sys("Creat write Error"); //now offset = 50 */
    return 0;
}
/* err_sys display error message and exit */
void err_sys(char *str)
{
    printf ("%s",str);
    exit (1);
}

```

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

31

pread() and pwrite() system call

```

#include <unistd.h>
ssize_t pread(int fd, void *buf, size_t nbyte, off_t offset);
// Returns: size of byte read, 0 if end of file, -1 on error
ssize_t pwrite(int fd, void *buf, size_t nbyte, off_t offset);
// Returns: size of byte write 0 if end of file, -1 on error

```

- Calling pread() (pwrite()) is equivalent to calling lseek() followed by a call to read() (write()) with following exceptions.

- There is no way to interrupt the two operations that occur when we call pread() (pwrite())
- The current file offset is not updated

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

32

```

//write.c shows example for pread() and pwrite()
#include <fcntl.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/stat.h>
#include <sys/types.h>

int main()
{
    int fd, nr, nr2, nw, nw2;
    char buf_w1[]="This is First message to write";
    char buf_w2[]="This is second message to write";
    char buf_rd[120];

    //open file
    fd = open("out.txt", O_RDWR | O_CREAT, 0666);
    // write first message to the file
    nw = pwrite(fd, buf_w1, strlen(buf_w1), 0);
    // write the second message just after the first message to the file
    nw2 = pwrite(fd, buf_w2, strlen(buf_w2), strlen(buf_w1)+1);
    // read the second message from the file
    nr = pread(fd, buf_rd, sizeof(buf_rd), strlen(buf_w1)+1);
    printf("read from file: %s",buf_rd);

    close(fd);

    return 0;
}

```

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

33

sync(), fsync() and fdatasync() System calls

- Traditional Unix system maintains a buffer cache (or page cache) in the kernel's space.
- When we write data to a file, the data is normally copied into the buffer cache and queued for writing to disk at some later time. (delayed write). The kernel eventually writes all the delayed-write blocks to disk, normally when it needs to reuse the buffer for some other disk block.
- sync(), fsync(), and fdatasync() system calls are provided to ensure consistence of file system on disk with the contents of buffer cache.
- The function sync() is normally called periodically (usually every 30 seconds) from a system daemon, often called update. This guarantees regular flushing of the kernel's block buffers.

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

34

sync(), fsync() and fdatasync() System calls

```

#include <unistd.h>
int fsync(int fd);
int fdatasync (int fd);
void sync (void);

```

Returns: 0 if ok, -1 on error

- The fsync() refers only to a single file, specified by the file descriptor filedes, and waits for the disk writes to complete before returning.
- This function is used when an application, such as a database, needs to be sure that the modified blocks have been written to the disk.
- The fdatasync() is similar to fsync(), but it affects only the data portions of a file. With fsync(), the file's attributes are also updated synchronously.

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

35

sync(), fsync() and fdatasync() System calls

```

#include <stdio.h>
#include <fcntl.h>

int main()
{
    char wbuffer[] = "1234567890";
    char rbuffer[100];
    int fd;

    /* Open the file */
    fd = open ("test.txt", O_RDWR | O_CREAT | O_TRUNC);
    /* Write 10 bytes of data */
    write (fd, (void *) wbuffer, 10);
    /* and make sure it's written */
    fsync (fd);
    /* Seek the beginning of the file */
    lseek (fd, 0, SEEK_SET);
    /* Read 10 bytes of data */
    read (fd, (void *) rbuffer, 10);
    /* Terminate the data we're read with a null character */
    rbuffer[10] = '\0';
    printf ("String read = %s.\n", rbuffer);
    close (fd);
    return 0;
}

```

COSC350 System Software, Fall 2024
Dr. Sang-Eon Park

36